

Document Sample Document

Page 1: Introduction to Retrieval-Augmented Generation

Retrieval-Augmented Generation (RAG) combines a retrieval system with a large language model. Instead of relying only on the model's trained knowledge, RAG fetches relevant text from an external knowledge base -- such as your uploaded PDFs -- and includes it in the prompt so the model can generate grounded, accurate answers.

Page 2: How This Sample Document Works

This file exists so you can test Document immediately after installation, without needing your own PDF. Try asking:

'What is Retrieval-Augmented Generation?'

'How does semantic search work in this app?'

The application will chunk this text, embed it, store the vectors in FAISS, and retrieve the most relevant passages to answer you.

Page 3: Semantic Search and FAISS

Semantic search finds text that is conceptually similar to a query, not just textually similar. Each chunk of text is converted into a high-dimensional vector (an embedding). FAISS indexes these vectors so that, given a new query vector, it can quickly find the nearest neighbors -- the chunks most likely to contain the answer.